

Music Generation With Machine Learning

Team 135: Romal Peccia, Abhishek Paul, Donald Lleshi
Supervisor: Jason Anderson Administrator: John Taglione

Background and Motivation


- Personal motivation: Writing music is hard!
- Desire for an application to extend user-recorded music
- ML is widely used for **random** generation of music
- ML for **input-based** generation of music is mostly in research phase, one commercial product was announced in Dec 2019

Goal

- Create a music composition application using neural networks.

Key Requirements

- **Input & Output Data:**
 - Data representation: (MIDI as Pianoroll)
 - Consistent duration: 5 seconds
 - Note range: 2 octaves
 - Musical Texture: Monophonic



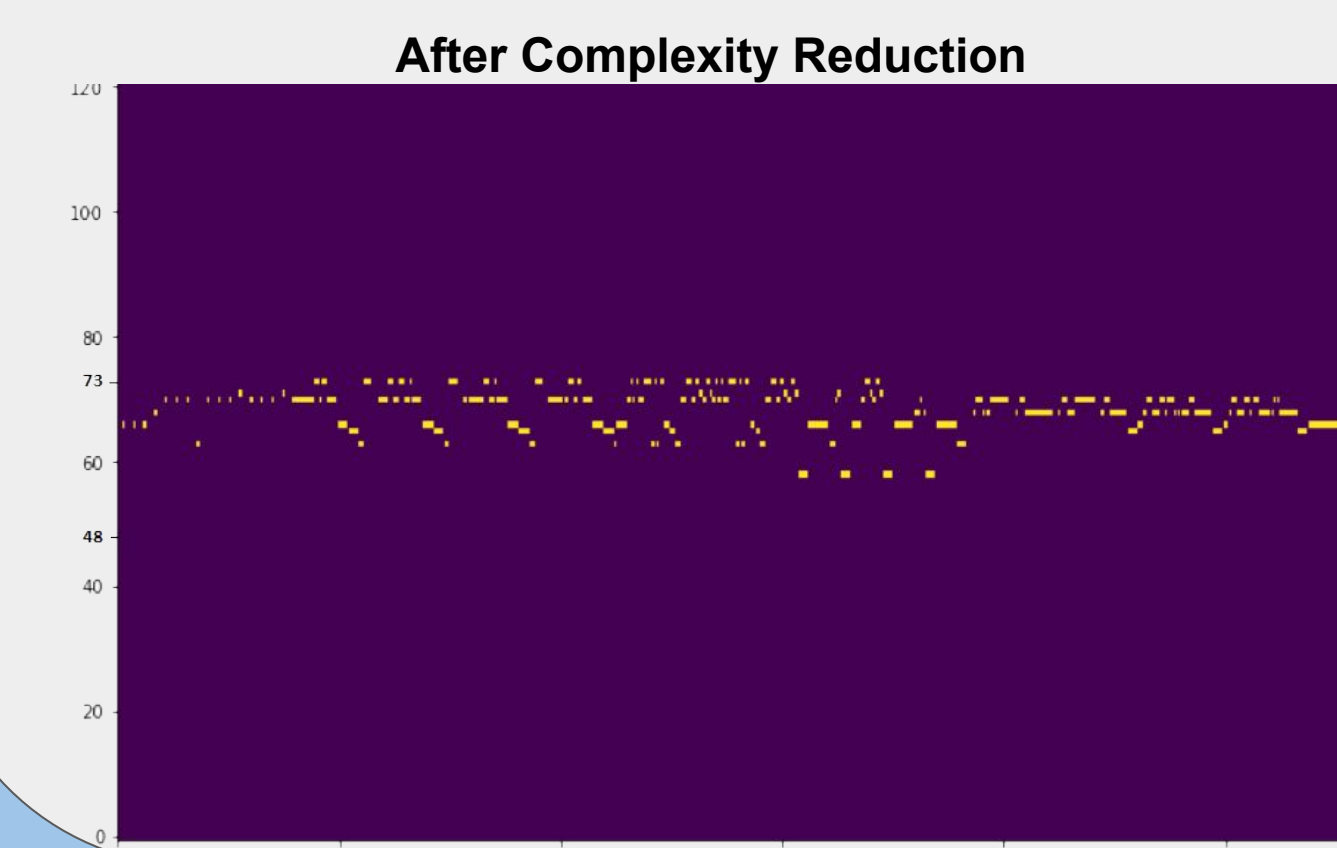
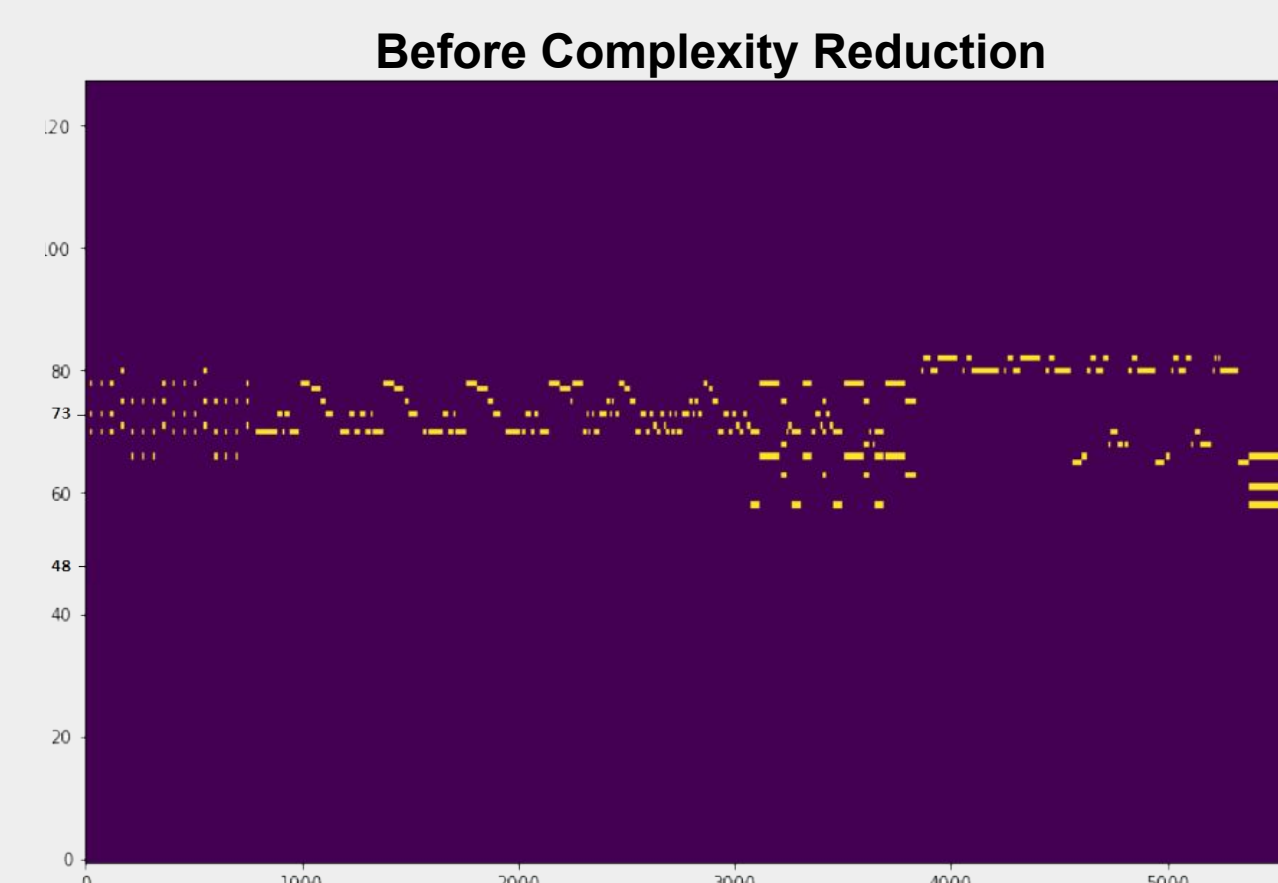
- **Model Architecture:**
 - Recurrent Neural Network (RNN)
- **User Interface:**
 - Interface between MIDI keyboard, computer, and RNN model
 - Return a 5 second extension of user's recorded 5 second melody

Design Challenges

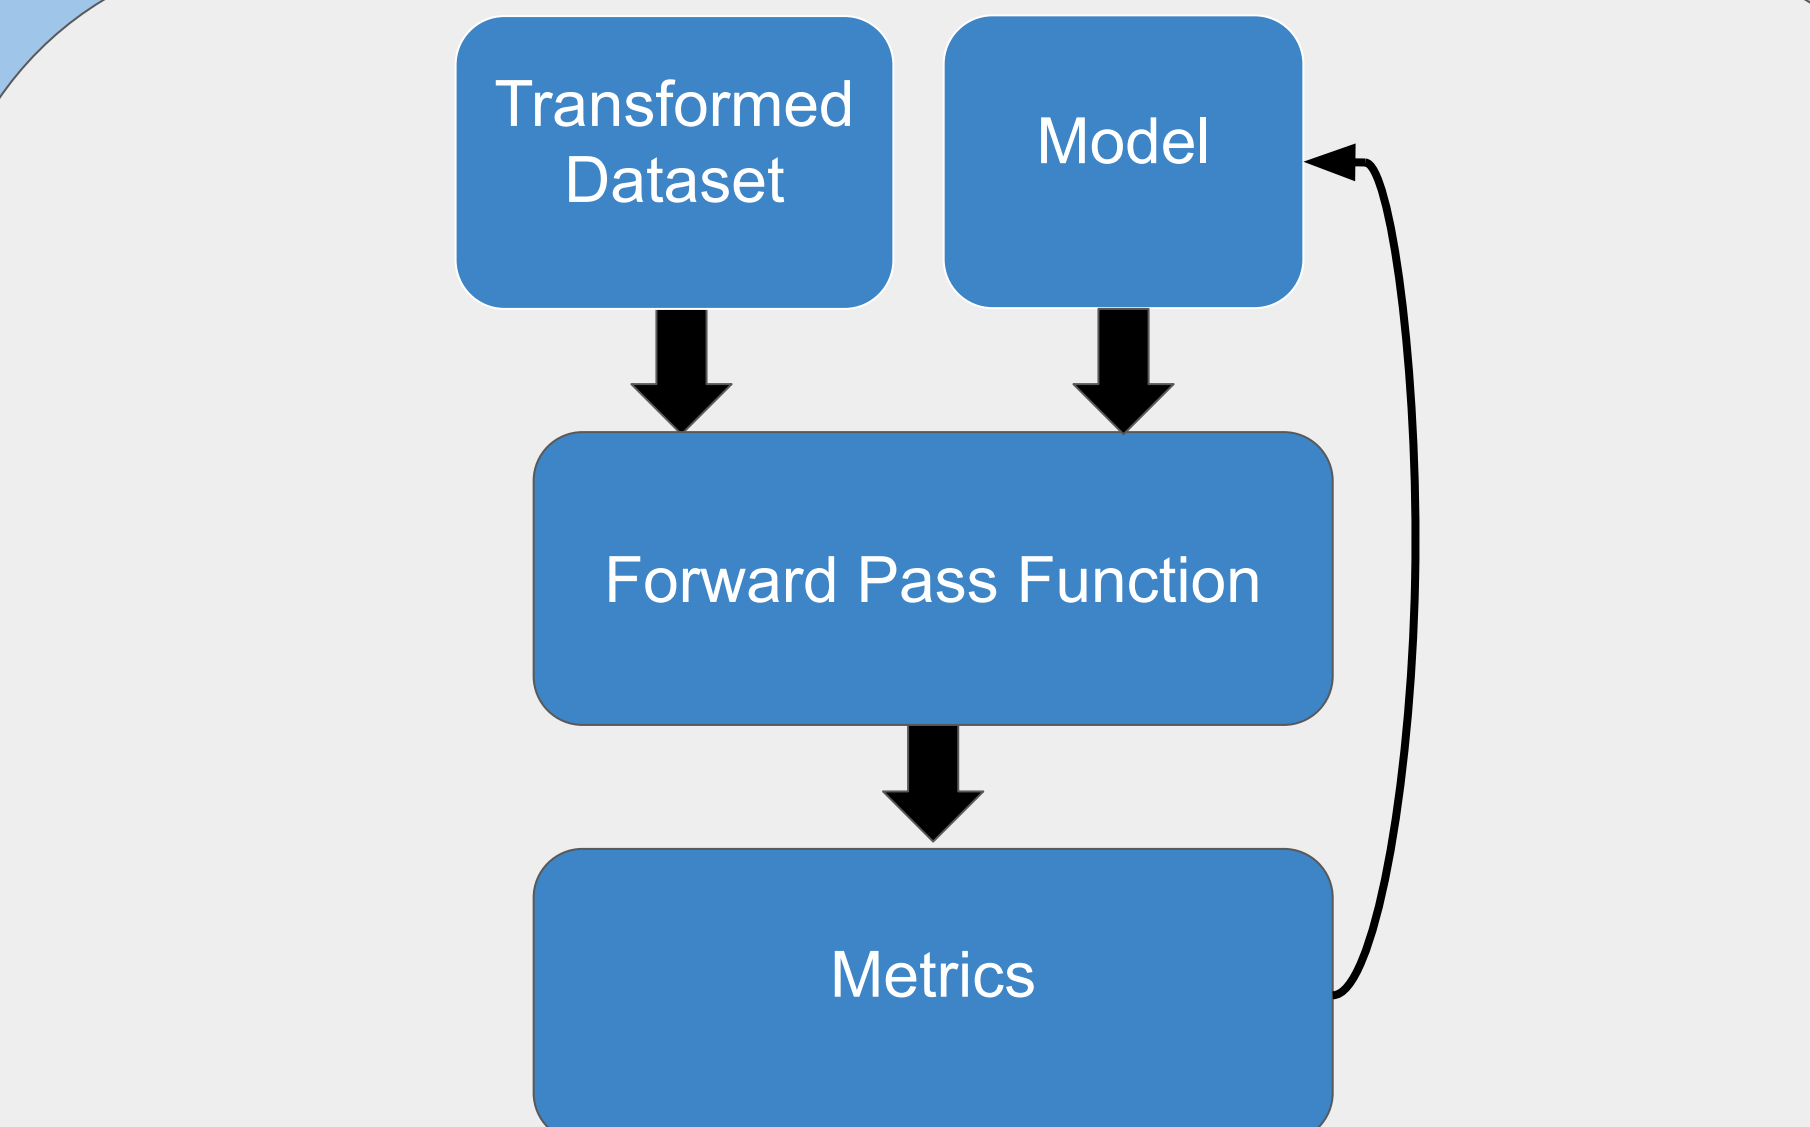
- **Data Representation of Music**
 - Typical ML tasks the team has experienced involve text, images, or spreadsheet data
 - MIDI in Pianoroll form was chosen after weighing alternatives
 - All functions that manipulated Pianoroll (other than read/write) were custom built by the team
- **Long model training times**
 - Complexity of dataset is 2^{128} combinations of notes per timestep (impossibly large)
 - Data transformations reduced complexity to 26 combinations (still days of training for a single model)
 - Reduction of scope: Focussed on training and testing one architecture (RNN) well rather than many with poorer results
- **Silent model outputs**
 - Music is more likely to be silent compared to playing any given note. Model generates silence in order to maximise accuracy
 - Custom accuracy metrics helped diagnose the problem
 - Improved data transformations were the solution

Data Processing

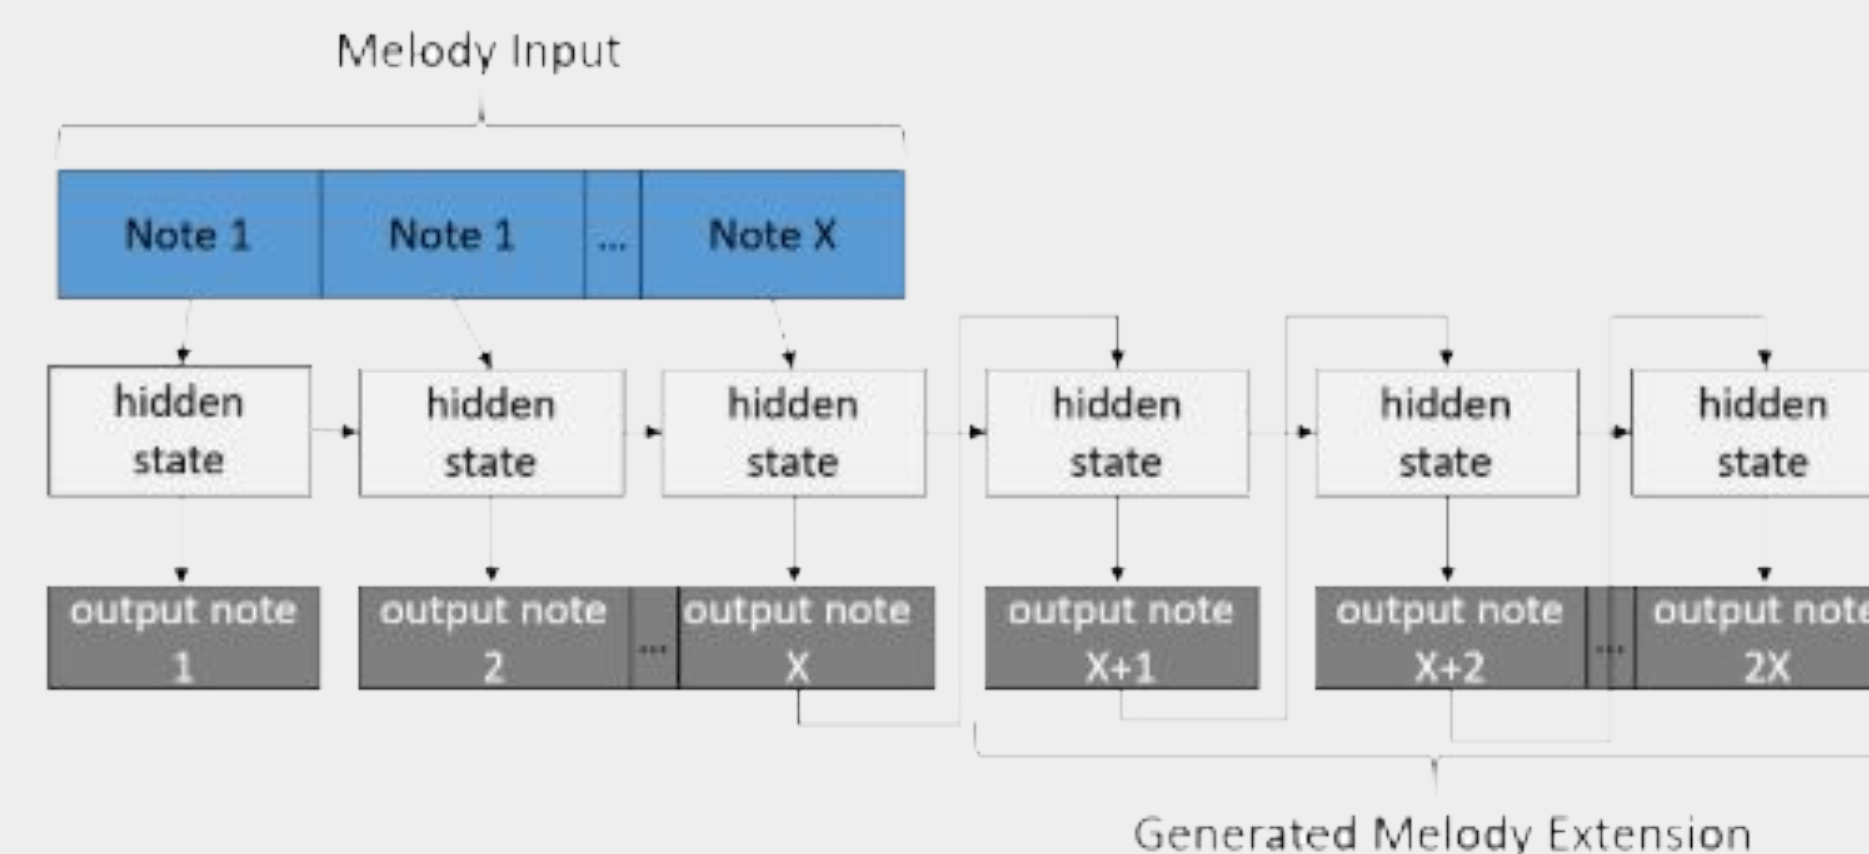
- Data representation: MIDI as Pianoroll
- **Acquired Dataset:**
 - 386 MIDI files (split into 6700) containing pop/classical music
- **Data Transformation**
 - Reducing Data Complexity
 - Polyphonic to monophonic
 - Note range reduction
 - Improving Data Quality
 - Split into 5s intervals
 - Tempo normalization
 - Silence threshold
 - Model Learnable Format
 - One-hot encoding
 - Melody/extension input/output pairs



Model Training



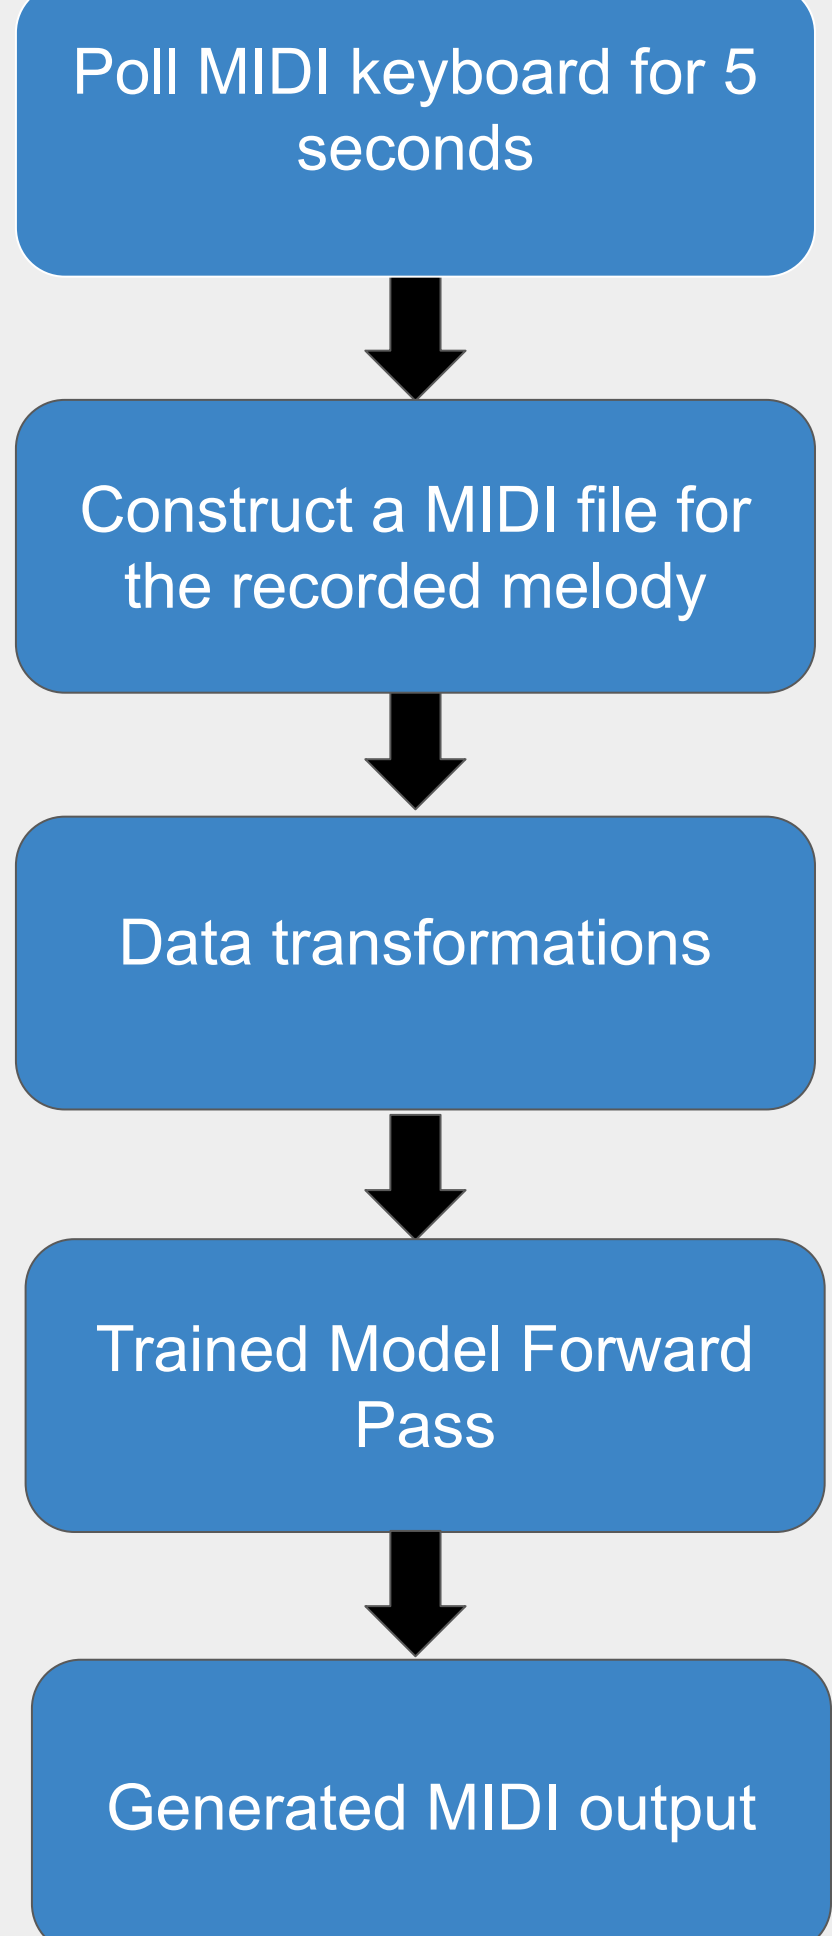
- **Model Architecture**
 - Gated Recurrent Unit Neural Network (RNN variant)
- **Forward Pass**
 - Melodies are passed through the model, extensions are generated through process below



- **Metrics**
 - Generated extensions are compared to true extensions to calculate accuracy and loss
 - Weights of model are iteratively updated using loss
 - Custom Zero and Non-Zero Accuracy (accuracy of predicting silence or a specific note respectively) metrics developed to further assess model biases

User Interface

- Framework built to:
 - Interface the MIDI keyboard with Python code
 - Read and interpret MIDI events
 - Construct a MIDI file (recorded melody)
- Data processing executed:
 - Data transformation functions
 - Model forward pass
- Model output converted to MIDI and played back



Results

- **Model Training Metrics**
 - Final model was trained for ~800 epochs (200 hours)
 - Accuracy metrics show how well model performed

Accuracy Type	Training	Validation
Zero	94.62 %	74.3572 %
Non-Zero	68.74294 %	45.6182 %
Overall	76.0272 %	48.6963%

- **Overall framework developed:**
 - Transformations of MIDI dataset to reduce complexity and pass through model
 - Model training process and custom metrics
 - Interface between keyboard, computer, and model
 - Generation of melody extensions based on user input in real time

Conclusions

- Accuracy metrics show that model was learning to generate music
- Although model can be improved, useful framework was developed to facilitate model learning and interfacing
- Future developers can mix and match datasets or models with our framework in order to find a better model, and plug their final model into our interface